

CS610 Assignment 2

Khushi Gupta

September 2025

Problem 1

Setup

- Implementations: **naive**, **input-blocked**, **output-blocked** 3D cross-correlation (no padding). Kernel: $3 \times 3 \times 3$.
- Input sizes tested: $N = \{128, 256, 512\}$ (i.e. $1 \ll 7, 1 \ll 8, 1 \ll 9$).
- Blocking configurable per-dimension (B_H, B_W, B_D). Each kernel timed 5 runs, reported time values are averages.
- PAPI events collected: PAPI_L1_DCM, PAPI_L2_DCM, PAPI_TOT_CYC.

Best block sizes

N	Block Sizes 1	Block Sizes 2
128	(4, 4, 64)	(1, 1, 64)
256	(1, 1, 128)	(1, 1, 256)
512	(1, 1, 256)	(1, 1, 512)

Table 1: Top two performing block sizes for each input size N .

Key Observed Trends

- **Input-blocked:** Slowest execution, with the fewest L1/L2 data cache misses but the highest cycle counts.
- **Output-blocked:** Fastest execution, with cycle counts and wall-clock time similar to the naïve version.

Possible reasons for input blocked to be slow -

1. **Many dirty evictions and stores:** Input-blocked repeatedly updates scattered output locations, causing cache lines to be reloaded and written back to memory multiple times. This leads to more evictions, extra stores, and higher cycles and time. In contrast, output-blocked writes to each cache line more contiguously, reducing these costs.
2. **Extra checks and indexing overhead:** input-blocked performs additional bound checks and index arithmetic (e.g., computing output indices and testing validity), increasing instruction count and branch overhead.

Results

```

Run 1
Naive Counters: L1 DCM=1040135, L2 DCM=2225319, CPU Cycles=61972692
Input Blocked: L1 DCM=1021989, L2 DCM=1374416, CPU Cycles=199004444
Output Blocked: L1 DCM=1099683, L2 DCM=1448103, CPU Cycles=64822656

Run 2
Naive Counters: L1 DCM=1041427, L2 DCM=2236517, CPU Cycles=62198608
Input Blocked: L1 DCM=1030505, L2 DCM=1379365, CPU Cycles=199951570
Output Blocked: L1 DCM=1087604, L2 DCM=1453411, CPU Cycles=63365114

Run 3
Naive Counters: L1 DCM=1038797, L2 DCM=2222812, CPU Cycles=61604289
Input Blocked: L1 DCM=1021459, L2 DCM=1378983, CPU Cycles=198061179
Output Blocked: L1 DCM=1091328, L2 DCM=1452413, CPU Cycles=63684440

Run 4
Naive Counters: L1 DCM=1040845, L2 DCM=2229413, CPU Cycles=62094533
Input Blocked: L1 DCM=1029037, L2 DCM=1379552, CPU Cycles=201865041
Output Blocked: L1 DCM=1099059, L2 DCM=1449344, CPU Cycles=64181221

Run 5
Naive Counters: L1 DCM=1038152, L2 DCM=2233796, CPU Cycles=60446471
Input Blocked: L1 DCM=1032957, L2 DCM=1372772, CPU Cycles=199958286
Output Blocked: L1 DCM=1108754, L2 DCM=1446069, CPU Cycles=65238036

=== Timing Results (averaged over 5 runs) ===
Naive : 23.0256 ms
Input Blocked : 66.3696 ms
Output Blocked : 23.5914 ms

```

(a) $N = 128$, Block = (4,4,64)

```

Run 1
Naive Counters: L1 DCM=1038217, L2 DCM=2232288, CPU Cycles=61556722
Input Blocked: L1 DCM=1038693, L2 DCM=2284011, CPU Cycles=209721570
Output Blocked: L1 DCM=1039354, L2 DCM=2160848, CPU Cycles=64952180

Run 2
Naive Counters: L1 DCM=1037462, L2 DCM=2162460, CPU Cycles=61719967
Input Blocked: L1 DCM=1022125, L2 DCM=2167493, CPU Cycles=207723167
Output Blocked: L1 DCM=1037131, L2 DCM=2166950, CPU Cycles=63365838

Run 3
Naive Counters: L1 DCM=1038302, L2 DCM=2164729, CPU Cycles=62106823
Input Blocked: L1 DCM=1025452, L2 DCM=2164941, CPU Cycles=208976474
Output Blocked: L1 DCM=1038641, L2 DCM=2168081, CPU Cycles=64416759

Run 4
Naive Counters: L1 DCM=1039246, L2 DCM=2164135, CPU Cycles=61980789
Input Blocked: L1 DCM=1025118, L2 DCM=2163451, CPU Cycles=210004177
Output Blocked: L1 DCM=1039941, L2 DCM=2170408, CPU Cycles=64870593

Run 5
Naive Counters: L1 DCM=1036759, L2 DCM=2171841, CPU Cycles=61085827
Input Blocked: L1 DCM=1024061, L2 DCM=2164131, CPU Cycles=209355118
Output Blocked: L1 DCM=1038536, L2 DCM=2171470, CPU Cycles=63628792

=== Timing Results (averaged over 5 runs) ===
Naive : 24.9344 ms
Input Blocked : 82.0566 ms
Output Blocked : 23.156 ms

```

(b) $N = 128$, Block = (1,1,64)

```

Run 1
Naive Counters: L1 DCM=10055148, L2 DCM=18782188, CPU Cycles=493233072
Input Blocked: L1 DCM=8685185, L2 DCM=18343632, CPU Cycles=1681603832
Output Blocked: L1 DCM=9862265, L2 DCM=18732815, CPU Cycles=499292026

Run 2
Naive Counters: L1 DCM=10127968, L2 DCM=18797068, CPU Cycles=494052610
Input Blocked: L1 DCM=8691065, L2 DCM=18362793, CPU Cycles=1680838163
Output Blocked: L1 DCM=9864683, L2 DCM=18788904, CPU Cycles=498470406

Run 3
Naive Counters: L1 DCM=10128717, L2 DCM=18810607, CPU Cycles=493353980
Input Blocked: L1 DCM=8677670, L2 DCM=18357342, CPU Cycles=1681677834
Output Blocked: L1 DCM=9872581, L2 DCM=18791736, CPU Cycles=498331609

Run 4
Naive Counters: L1 DCM=10089800, L2 DCM=18791149, CPU Cycles=492025327
Input Blocked: L1 DCM=8676893, L2 DCM=18313567, CPU Cycles=1682488877
Output Blocked: L1 DCM=9793146, L2 DCM=18718524, CPU Cycles=499203519

Run 5
Naive Counters: L1 DCM=10072017, L2 DCM=18740587, CPU Cycles=493214587
Input Blocked: L1 DCM=8679107, L2 DCM=18318417, CPU Cycles=1681274455
Output Blocked: L1 DCM=9617931, L2 DCM=18725177, CPU Cycles=499152172

=== Timing Results (averaged over 5 runs) ===
Naive : 174.138 ms
Input Blocked : 594.369 ms
Output Blocked : 159.135 ms

```

(c) $N = 256$, Block = (1,1,256)

```

Run 1
Naive Counters: L1 DCM=10004631, L2 DCM=18514015, CPU Cycles=495092078
Input Blocked: L1 DCM=9533217, L2 DCM=18547138, CPU Cycles=1540114363
Output Blocked: L1 DCM=9739104, L2 DCM=18589716, CPU Cycles=505446426

Run 2
Naive Counters: L1 DCM=9979171, L2 DCM=18627493, CPU Cycles=494029689
Input Blocked: L1 DCM=9531488, L2 DCM=18575600, CPU Cycles=1538023915
Output Blocked: L1 DCM=9750963, L2 DCM=18675099, CPU Cycles=504265558

Run 3
Naive Counters: L1 DCM=9988868, L2 DCM=18621892, CPU Cycles=494990266
Input Blocked: L1 DCM=9526090, L2 DCM=18543659, CPU Cycles=1538048924
Output Blocked: L1 DCM=9759080, L2 DCM=18687306, CPU Cycles=503484594

Run 4
Naive Counters: L1 DCM=9999333, L2 DCM=18608926, CPU Cycles=493903960
Input Blocked: L1 DCM=9532526, L2 DCM=18567210, CPU Cycles=1537714296
Output Blocked: L1 DCM=9751737, L2 DCM=18574721, CPU Cycles=505111768

Run 5
Naive Counters: L1 DCM=9975329, L2 DCM=18561692, CPU Cycles=494324159
Input Blocked: L1 DCM=9533617, L2 DCM=18583488, CPU Cycles=1538505247
Output Blocked: L1 DCM=9761621, L2 DCM=18640070, CPU Cycles=505515538

=== Timing Results (averaged over 5 runs) ===
Naive : 153.999 ms
Input Blocked : 564.778 ms
Output Blocked : 146.417 ms

```

(d) $N = 256$, Block = (1,1,128)

```

Run 1
Naive Counters: L1 DCM=363486505, L2 DCM=148975675, CPU Cycles=4797733980
Input Blocked: L1 DCM=161110391, L2 DCM=147829819, CPU Cycles=10193298602
Output Blocked: L1 DCM=341258427, L2 DCM=152628224, CPU Cycles=3664664205

Run 2
Naive Counters: L1 DCM=333615507, L2 DCM=151958745, CPU Cycles=3690936586
Input Blocked: L1 DCM=157168996, L2 DCM=149317818, CPU Cycles=9140007391
Output Blocked: L1 DCM=341772254, L2 DCM=153026261, CPU Cycles=3660284707

Run 3
Naive Counters: L1 DCM=333055183, L2 DCM=151043976, CPU Cycles=3682241411
Input Blocked: L1 DCM=156561975, L2 DCM=148341371, CPU Cycles=9235542240
Output Blocked: L1 DCM=34091946, L2 DCM=152131567, CPU Cycles=3671892814

Run 4
Naive Counters: L1 DCM=332644603, L2 DCM=152172278, CPU Cycles=3665116513
Input Blocked: L1 DCM=158011288, L2 DCM=149275406, CPU Cycles=9145798370
Output Blocked: L1 DCM=339271460, L2 DCM=153945272, CPU Cycles=3640499862

Run 5
Naive Counters: L1 DCM=331288713, L2 DCM=151439711, CPU Cycles=3685533240
Input Blocked: L1 DCM=158076331, L2 DCM=149283973, CPU Cycles=9137406195
Output Blocked: L1 DCM=337966450, L2 DCM=152797618, CPU Cycles=3661923822

=== Timing Results (averaged over 5 runs) ===
Naive : 969.785 ms
Input Blocked : 2298.24 ms
Output Blocked : 913.453 ms

```

(e) $N = 512$, Block = (1,1,512)

```

Run 1
Naive Counters: L1 DCM=333497784, L2 DCM=152294598, CPU Cycles=3698261543
Input Blocked: L1 DCM=161857138, L2 DCM=150800877, CPU Cycles=11399322230
Output Blocked: L1 DCM=344352449, L2 DCM=153256301, CPU Cycles=3697725651

Run 2
Naive Counters: L1 DCM=336069899, L2 DCM=150665295, CPU Cycles=3723204279
Input Blocked: L1 DCM=161706068, L2 DCM=149900631, CPU Cycles=11395860149
Output Blocked: L1 DCM=343849619, L2 DCM=152613608, CPU Cycles=3712794856

Run 3
Naive Counters: L1 DCM=337868456, L2 DCM=152154101, CPU Cycles=3739269755
Input Blocked: L1 DCM=161655973, L2 DCM=149936465, CPU Cycles=11395900480
Output Blocked: L1 DCM=343840127, L2 DCM=152907029, CPU Cycles=3695827906

Run 4
Naive Counters: L1 DCM=336963238, L2 DCM=152324914, CPU Cycles=3699605198
Input Blocked: L1 DCM=161607723, L2 DCM=150758201, CPU Cycles=11396624506
Output Blocked: L1 DCM=343824529, L2 DCM=153041165, CPU Cycles=3701339823

Run 5
Naive Counters: L1 DCM=337132304, L2 DCM=152401482, CPU Cycles=3698879226
Input Blocked: L1 DCM=161558264, L2 DCM=151184063, CPU Cycles=11394478687
Output Blocked: L1 DCM=343958707, L2 DCM=152801749, CPU Cycles=3707370355

=== Timing Results (averaged over 5 runs) ===
Naive : 828.833 ms
Input Blocked : 2510.22 ms
Output Blocked : 815.854 ms

```

(f) $N = 512$, Block = (1,1,256)

Figure 1: Best two block configurations per input size N . Each subfigure shows the performance for one block size.

Problem 2

N threads read N files and count words & lines. The original code performed updates to shared counters inside inner loops, causing heavy lock contention and cache-coherence (false sharing) traffic.

Key performance issues in original code

- Per-word `pthread_mutex_lock/unlock` on `tracker.total_words_processed`.
- Frequent writes to contiguous `tracker.word_count[]` entries from different threads - false sharing and many HITM cacheline transfers.
- Per-line locking for line counter (lock/unlock per line).

Optimizations applied

- Local aggregation of counters inside each thread (`word_count`, `line_count`).
- Single write to the per-thread slot (`tracker.word_count[tid]`) and one short critical section to add local totals to global counters at thread end.
- Removed per-word and per-line frequent locking.

Observed effect

- **Before:** Shared cachelines showed many HITM events indicating frequent ownership transfers between cores. These events mean that one core attempted to read or write a cacheline owned by another core, triggering coherence traffic and stalls. The two snapshots below illustrate the high level of coherence activity and contention on the shared data cachelines:

Shared Data Cache Line Table (27 entries, sorted on Total HITMs)																
Index	Address	Node	PA	cnt	Tot Hitm	Total	Load LclHitm	Hitm	RmtHitm	Total records	Total Loads	Total Stores	L1Hit	L1Miss	N/A	Cor
0	0x5711606f93c0	0	832	36.08%	105	105	0	1231	1041	190	188	2	0	546		
1	0xfffff8b6d818ae080	0	221	35.05%	102	102	0	377	308	69	68	1	0	46		
2	0x5711606f93c0	0	1	12.37%	36	36	0	45	45	0	0	0	0	0		
3	0x5711606f9380	0	86	6.87%	20	20	0	103	69	34	2	32	0	2		
4	0xfffff8b70cc321800	0	1	1.03%	3	3	0	3	3	0	0	0	0	0		
5	0xfffff8b6d95d82980	0	1	0.69%	2	2	0	3	3	0	0	0	0	0		
6	0xfffff8b6e28883280	0	1	0.69%	2	2	0	2	2	0	0	0	0	0		

Figure 2: Before optimization - cacheline coherence records

Shared Data Cache Line Table (22 entries, sorted on Total HITMs)																
Index	Address	Node	PA	cnt	Tot Hitm	Total	Load LclHitm	Hitm	RmtHitm	Total records	Total Loads	Total Stores	L1Hit	L1Miss	N/A	Cor
0	0xfffff8b6d818a0880	0	267	42.62%	130	130	0	443	373	70	70	0	0	2		
1	0x6324259c93c0	0	784	32.13%	98	98	0	1144	978	166	164	2	0	0		
2	0x6324259c93c0	0	1	12.46%	38	38	0	45	45	0	0	0	0	0		
3	0x6324259c9380	0	75	5.90%	18	18	0	93	68	25	0	25	0	0		
4	0xfffff8b6f28f48940	0	1	0.66%	2	2	0	2	2	0	0	0	0	0		
5	0xfffff8b70cc2b5600	0	2	0.66%	2	2	0	3	3	0	0	0	0	0		
6	0xfffff8b70cc2b5a00	0	1	0.66%	2	2	0	2	2	0	0	0	0	0		

Figure 3: Before optimization - cacheline coherence records

- **After:** The Shared Data Cache Line Table shows that inter-core cacheline transfers were almost entirely eliminated. In the optimized run, the previously affected cachelines no longer

appeared as sources of contention. Only a single residual entry was recorded with just one HITM event. This negligible activity arises because different threads still update their own per-thread word counters, which may occasionally map to the same cacheline :

Shared	Data	Cache	Line	Table	(0 entries, sorted on Total HITMs)																
Index			Cacheline	Address	Node	PA	cnt	Tot	Load	Hitm		Total	LclHitm	RmtHitm	Total	records	Total	Stores			
								Hitm										L1Hit	L1Miss	N/A	

Figure 4: After optimization - cacheline coherence records.

Shared Data Cache Line Table (1 entries, sorted on Total HITMs)																		
Index	Cacheline	Address	Node	PA	cnt	Tot Hitm	Total	Load Hitm LcLHitm	RmtHitm	Total records	Total Loads	Total Stores	Stores L1Hit	Stores L1Miss	N/A	Core FB	Load L1	Hit
0	0xfffff8b6d80ee8000		0	1	100.00%	1	1	1	0	1	1	0	0	0	0	0	0	0

Figure 5: After optimization - cacheline coherence records.

Problem 3

Results

The following results were obtained on *image1.cse.iitk.ac.in*. For 64 threads the execution could not complete in reasonable time.

Lock / # threads	1	2	4	8	16	32
Pthread mutex	Run1: 30 Run2: 31 Run3: 28 Run4: 24 Run5: 32 Avg: 29	Run1: 425 Run2: 262 Run3: 517 Run4: 379 Run5: 444 Avg: 405	Run1: 1536 Run2: 1121 Run3: 1399 Run4: 1113 Run5: 1860 Avg: 1405	Run1: 8245 Run2: 8460 Run3: 6961 Run4: 6904 Run5: 7014 Avg: 7516	Run1: 45940 Run2: 43667 Run3: 42118 Run4: 36090 Run5: 42287 Avg: 42020	Run1: 210303 Run2: 179210 Run3: 175000 Run4: 172615 Run5: 178313 Avg: 183088
Filter lock	Run1: 9 Run2: 5 Run3: 5 Run4: 5 Run5: 5 Avg: 6	Run1: 722 Run2: 554 Run3: 810 Run4: 678 Run5: 722 Avg: 697	Run1: 4597 Run2: 4087 Run3: 4395 Run4: 4304 Run5: 4590 Avg: 4395	Run1: 119103 Run2: 122311 Run3: 115104 Run4: 139701 Run5: 126059 Avg: 124456	Run1: 1360747 Run2: 1351788 Run3: 1351275 Run4: 1362924 Run5: 1372719 Avg: 1359891	Run1: 11245362 Run2: 11025148 Run3: 11190998 Run4: 11167281 Run5: 11132490 Avg: 11152256
Bakery lock	Run1: 74 Run2: 48 Run3: 48 Run4: 64 Run5: 46 Avg: 56	Run1: 864 Run2: 720 Run3: 904 Run4: 946 Run5: 852 Avg: 857.2	Run1: 3760 Run2: 3600 Run3: 11148 Run4: 3788 Run5: 4576 Avg: 5374	Run1: 44536 Run2: 18384 Run3: 44568 Run4: 34160 Run5: 34000 Avg: 35130	Run1: 252431 Run2: 253678 Run3: 245631 Run4: 250302 Run5: 250706 Avg: 250550	Run1: 1384240 Run2: 1369689 Run3: 1406609 Run4: 1455927 Run5: 1388586 Avg: 1401010
Spin lock	Run1: 18 Run2: 11 Run3: 13 Run4: 24 Run5: 11 Avg: 15	Run1: 383 Run2: 390 Run3: 767 Run4: 286 Run5: 326 Avg: 430	Run1: 3474 Run2: 2115 Run3: 2039 Run4: 2557 Run5: 2895 Avg: 2616	Run1: 15129 Run2: 17856 Run3: 15239 Run4: 15393 Run5: 17702 Avg: 16264	Run1: 111595 Run2: 101857 Run3: 120797 Run4: 112310 Run5: 107725 Avg: 110857	Run1: 687443 Run2: 729216 Run3: 868888 Run4: 761947 Run5: 694699 Avg: 748439
Ticket lock	Run1: 31 Run2: 44 Run3: 39 Run4: 68 Run5: 40 Avg: 44	Run1: 408 Run2: 406 Run3: 776 Run4: 488 Run5: 560 Avg: 528	Run1: 6883 Run2: 2428 Run3: 2640 Run4: 6503 Run5: 2408 Avg: 4172	Run1: 25816 Run2: 27266 Run3: 25544 Run4: 12692 Run5: 27084 Avg: 23680	Run1: 111918 Run2: 121398 Run3: 119586 Run4: 121899 Run5: 65724 Avg: 108105	Run1: 507119 Run2: 521938 Run3: 508877 Run4: 498921 Run5: 516483 Avg: 510668
Array Q lock	Run1: 11 Run2: 24 Run3: 13 Run4: 22 Run5: 15 Avg: 17	Run1: 548 Run2: 618 Run3: 1878 Run4: 948 Run5: 748 Avg: 948	Run1: 3116 Run2: 2592 Run3: 2692 Run4: 2704 Run5: 2840 Avg: 2789	Run1: 23199 Run2: 22811 Run3: 18767 Run4: 28134 Run5: 12776 Avg: 21137	Run1: 50287 Run2: 69180 Run3: 66194 Run4: 73406 Run5: 69052 Avg: 65624	Run1: 549321 Run2: 543002 Run3: 561290 Run4: 541796 Run5: 545946 Avg: 548271

Table 2: All these were recorded for input size $N = 1e6$

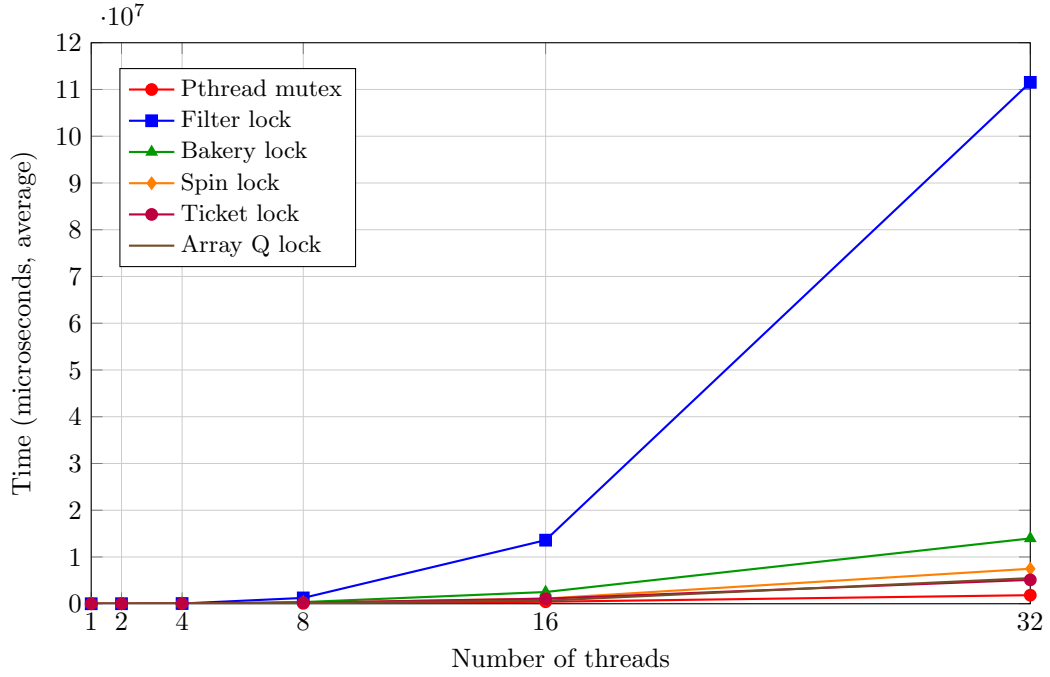


Figure 6: Average time (μs) vs number of threads for each lock type.

Optimizations

- Use of `pause` instruction inside spin-waits to reduce pipeline contention and busy-loop aggressiveness.
- Explicit `mfence` calls where needed to enforce ordering between reads/writes (Filter, Bakery, Ticket).
- Array-based queue uses cache-line alignment/padding (`alignas(64)/padding`) so each waiter spins on its own cache line to avoid sharing.

Trends

- At high thread counts, FilterLock and BakeryLock degrade noticeably due to their $O(N)$ scanning/booking behavior.
- TicketLock and ArrayQLock provide better stability. TicketLock gives orderly FIFO handoff, while ArrayQLock reduces coherence by having each waiter spin on a separate, padded slot. Both perform better compared with a spinlock at high number of threads.
- Adding `pause` instruction in loops and cache-line padding for ArrayQLock improved performance.